

Homework2

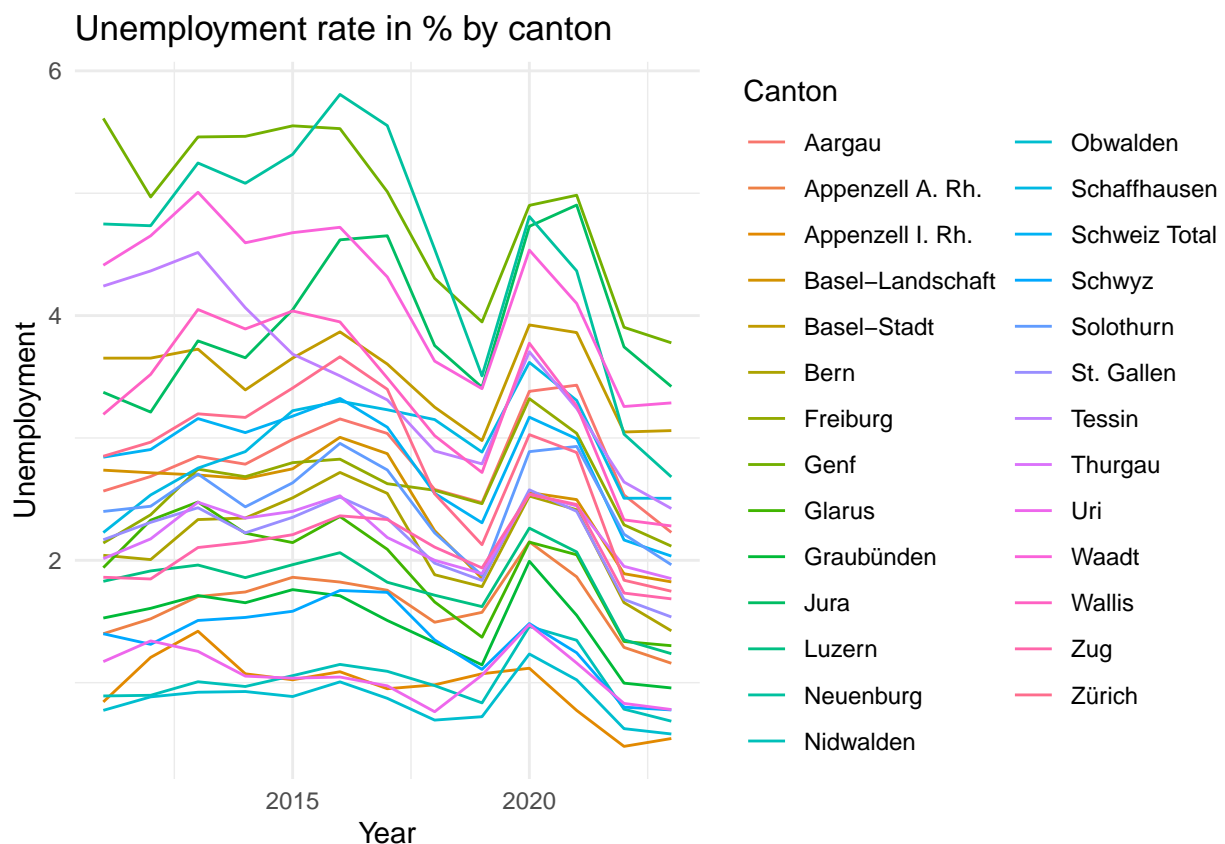
Erdan Beka, Cyril Scheuermann, Roman Krass, Keijo Nierula

2024-05-21

Data exploration

In this part we visualize the raw data to have an overview of the unemployment rate. We will do this for each canton and then summarize the data to get an overview of the unemployment rate in Switzerland.

```
# visualize the data for each canton
data |>
  ggplot(aes(x = Year, y = Unemployment, color = Canton)) +
  geom_line() +
  theme_minimal() +
  labs(title = "Unemployment rate in % by canton", color = "Canton")
```




In this plot we can see the unemployment rate for each canton in % over the years. We can see that the

percentage of unemployment is different for each canton but they all follow the same trend. This is for example visible in the year 2020 where the unemployment rate increased for all cantons.

```
# visualize the data summarized
data |>
  index_by(Year) |>
  summarise(Unemployment = sum(Unemployment_number, na.rm = FALSE)) |>
  ggplot(aes(x = Year, y = Unemployment)) +
  geom_line() +
  theme_minimal() +
  labs(title = "Unemployment rate absolute")
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').
```



In this plot we can see the absolute number of unemployed people in Switzerland over the years. We can see that the number of unemployed people increased in 2020 and decreased a lot in 2022.

Forecasting

Liner Trend Model


```

# Define trend model for the Unemployment
fit <- tslm(Unemployment ~ trend, data = data_ts)

# Forecast the unemployment rate for the next 5 years
fit_fc <- fit |>
  forecast(h = 5)

# Plot the forecast
fit_fc |>
  autoplot()

```



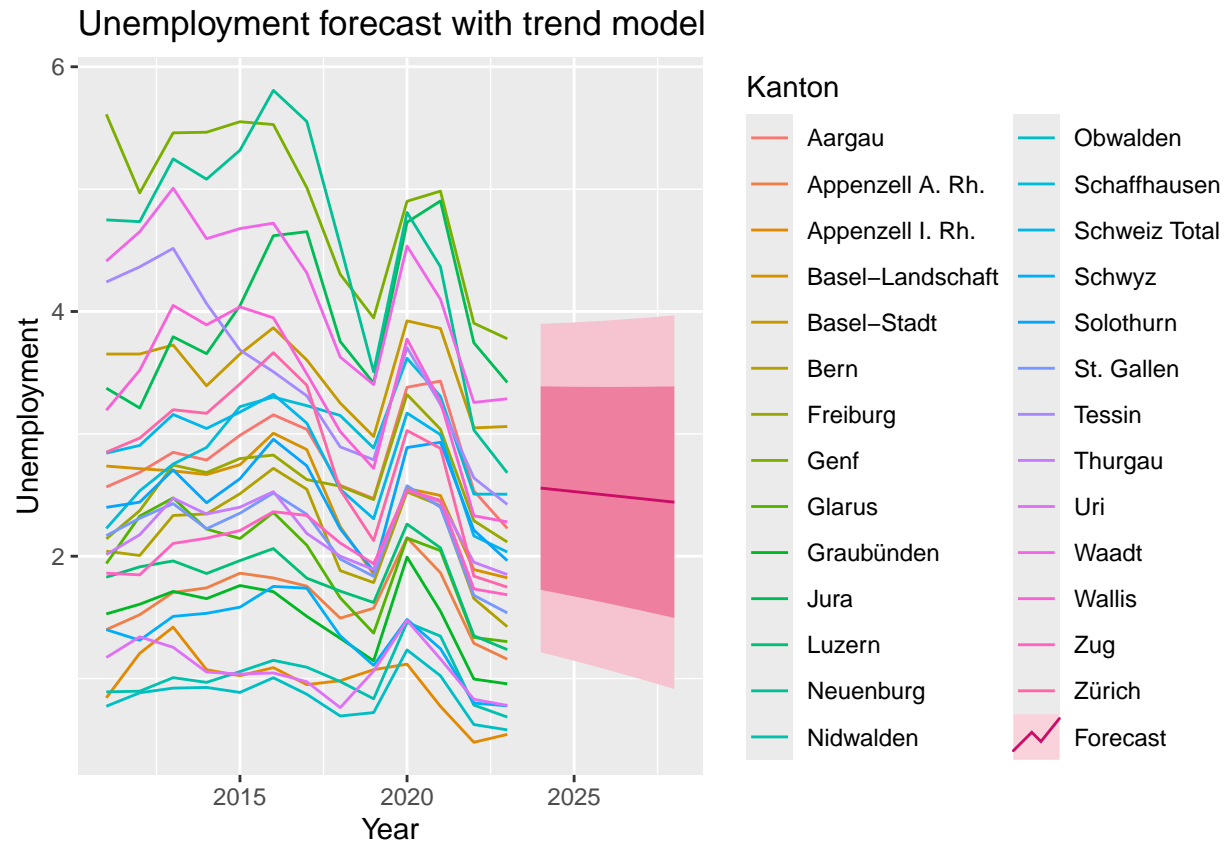
When we look at the linear trend model we can see that the unemployment rate is forecasted to decrease over the next 5 years. The range of the forecast is quite large. But this is also because of the big decrease in overall unemployment in 2021 and 2022.

visualize Results

```

# Plot forecast for all cantons
autoplot(data_tsibble) +
  autolayer(fit_fc, series = "Forecast") +
  xlab("Year") +
  ylab("Unemployment") +
  ggtitle("Unemployment forecast with trend model")

```



When we look at the forecast for all cantons we can see that the unemployment rate is forecasted to decrease. But because the forecast is combined for all cantons we can't see the differences between the cantons. Because the unemployment rate is combined from all cantons the range isn't suitable for all cantons. For example Appenzell Innerrhoden has a very low unemployment rate compared to other cantons. The forecast is why higher than the current unemployment rate. Which is not what the general forecast shows.

Next we try the mean method to get another forecast:

```
# Split the data into training and test sets
train_data <- data_tsibble |>
  filter(Year <= 2018)

test_data <- data_tsibble |>
  filter(Year > 2018)

# Mean Forecasting Method
mean_model <- train_data |>
  model(mean_fc = MEAN(Unemployment))

mean_fc <- mean_model |>
  forecast(new_data = test_data)

# Plot the forecast
autoplot(train_data, series = "Training Data") +
  autolayer(test_data, series = "Test Data") +
```

```

autolayer(mean_fc, series = "Mean Forecast") +
xlab("Year") +
ylab("Unemployment") +
ggtitle("Unemployment Forecast with Mean Method")

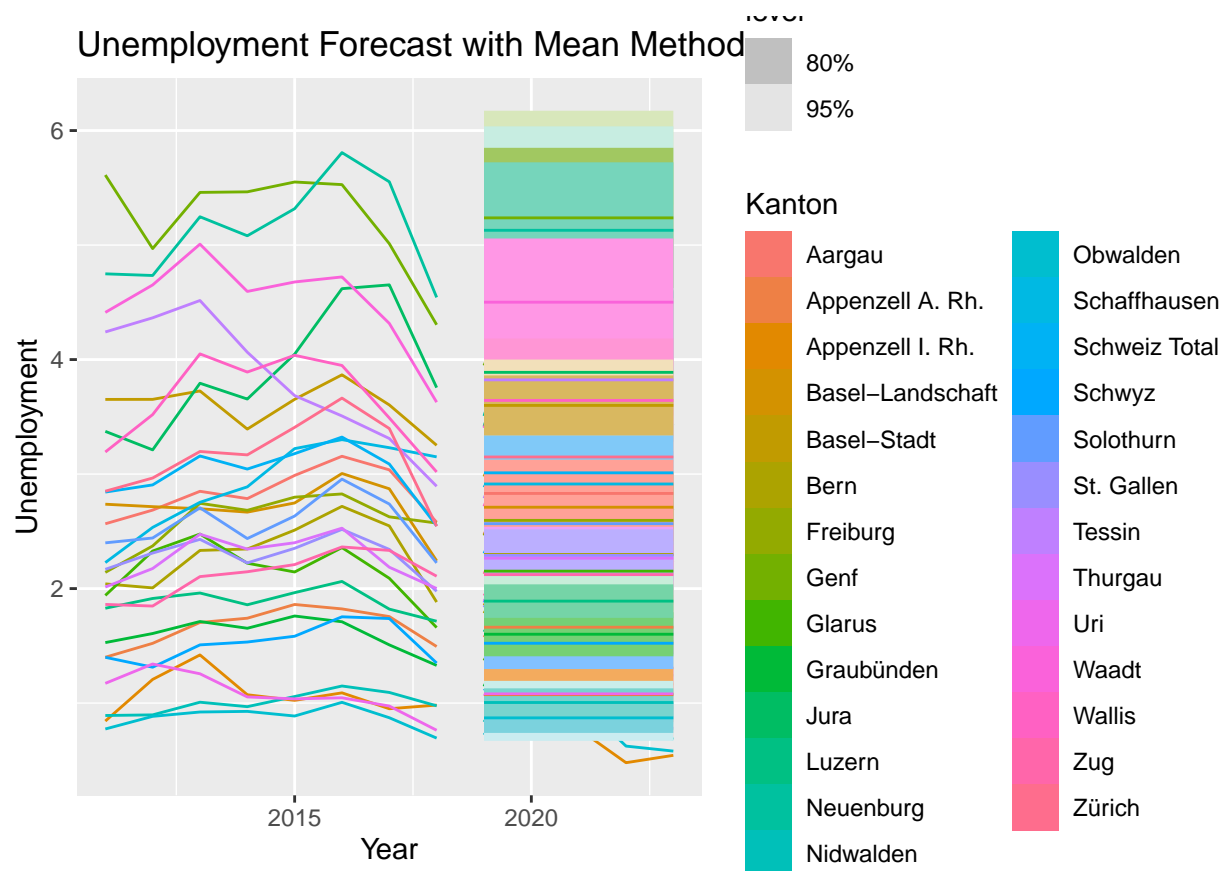
```

```
## Warning in geom_line(...): Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(eval_tidy(expr(aes(!!!aes_spec)))), data = object, ..., :
## Ignoring unknown parameters: 'series'
```

```
## Warning in ggdist::geom_lineribbon(without(intvl_mapping, "colour_ramp"), :
## Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(mapping = without(mapping, "shape"), data =
## unpack_data(object[single_row[["FALSE"]], : Ignoring unknown parameters:
## 'series'
```



Naive method:

```

# Naive Forecasting Method
naive_model <- train_data |>
  model(naive_fc = NAIVE(Unemployment))

naive_fc <- naive_model |>

```

```

forecast(new_data = test_data)

# Plot the forecast
autoplot(train_data, series = "Training Data") +
  autolayer(test_data, series = "Test Data") +
  autolayer(naive_fc, series = "Naïve Forecast") +
  xlab("Year") +
  ylab("Unemployment") +
  ggtitle("Unemployment Forecast with Naïve Method")


```

```
## Warning in geom_line(...): Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(eval_tidy(expr(aes(!!!aes_spec)))), data = object, ..., :
## Ignoring unknown parameters: 'series'
```

```
## Warning in ggdist::geom_lineribbon(without(intvl_mapping, "colour_ramp"), :
## Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(mapping = without(mapping, "shape"), data =
## unpack_data(object[single_row[["FALSE"]], : Ignoring unknown parameters:
## 'series'
```



Drift Method:

```
# Drift Forecasting Method
drift_model <- train_data |>
  model(drift_fc = RW(Unemployment ~ drift()))

drift_fc <- drift_model |>
  forecast(new_data = test_data)


# Plot the forecast
autoplot(train_data, series = "Training Data") +
  autolayer(test_data, series = "Test Data") +
  autolayer(drift_fc, series = "Drift Forecast") +
  xlab("Year") +
  ylab("Unemployment") +
  ggtitle("Unemployment Forecast with Drift Method")
```

```
## Warning in geom_line(...): Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(eval_tidy(expr(aes(!!!aes_spec)))), data = object, ..., :
## Ignoring unknown parameters: 'series'
```

```
## Warning in ggdist::geom_lineribbon(without(intvl_mapping, "colour_ramp"), :
## Ignoring unknown parameters: 'series'
```

```
## Warning in geom_line(mapping = without(mapping, "shape"), data =
## unpack_data(object[single_row[["FALSE"]]), : Ignoring unknown parameters:
## 'series'
```



Just from looking at the plots, it is hard to tell which method is the best. We will now calculate the accuracy of the forecasts to get a better idea of which method is the best.

```
# Calculate accuracy metrics
mean_accuracy <- mean_fc |>
  accuracy(test_data)

naive_accuracy <- naive_fc |>
  accuracy(test_data)

drift_accuracy <- drift_fc |>
  accuracy(test_data)

mean_RMSE <- mean(mean_accuracy$RMSE)
naive_RMSE <- mean(naive_accuracy$RMSE)
drift_RMSE <- mean(drift_accuracy$RMSE)
cat("Mean RMSE: ", mean_RMSE, "\n")
```

```
## Mean RMSE: 0.6047311
```

```
cat("Naïve RMSE: ", naive_RMSE, "\n")
```

```
## Naïve RMSE: 0.4651364
```

```
cat("Drift RMSE: ", drift_RMSE, "\n")
```

```
## Drift RMSE: 0.5048573
```

```
cat("We can see that the Naïve method has the lowest RMSE, which means it is the most accurate method.")
```

```
## We can see that the Naïve method has the lowest RMSE, which means it is the most accurate method.
```

```
mean(naive_accuracy$MAPE)
```

```
## [1] 20.45493
```

We can see that the Naïve method has the lowest RMSE, which means it is the most accurate method. It also has the lowest MAPE (20.45), which means it is the most accurate method so far.

To get a better idea of the accuracy we visualize the performance of the models by combining the residuals:

```
# Residuals for each model
mean_resid <- augment(mean_model) |>
  mutate(model = "Mean")

naive_resid <- augment(naive_model) |>
  mutate(model = "Naïve")

drift_resid <- augment(drift_model) |>
  mutate(model = "Drift")
```



```

# Combine residuals
residuals <- bind_rows(mean_resid, naive_resid, drift_resid)


# Plot residuals
residuals |>
  ggplot(aes(x = Year, y = .resid, color = model)) +
  geom_line() +
  facet_wrap(~model, scales = "free_y") +
  theme_minimal() +
  labs(title = "Residuals of Forecasting Models", y = "Residuals")

```

```

## Warning: Removed 54 rows containing missing values or values outside the scale range
## ('geom_line()').

```



```

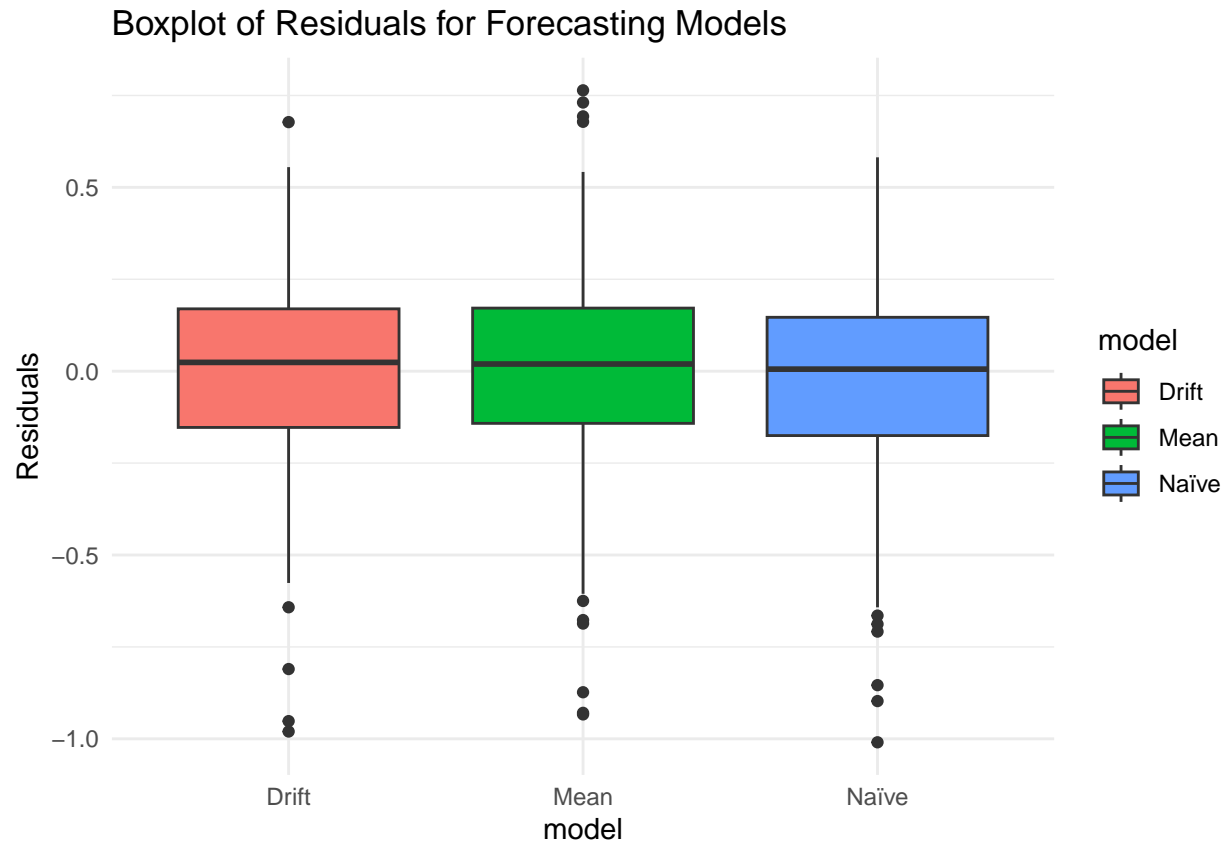
# Boxplot of residuals
residuals |>
  ggplot(aes(x = model, y = .resid, fill = model)) +
  geom_boxplot() +
  theme_minimal() +
  labs(title = "Boxplot of Residuals for Forecasting Models", y = "Residuals")

```

```

## Warning: Removed 54 rows containing non-finite outside the scale range
## ('stat_boxplot()').

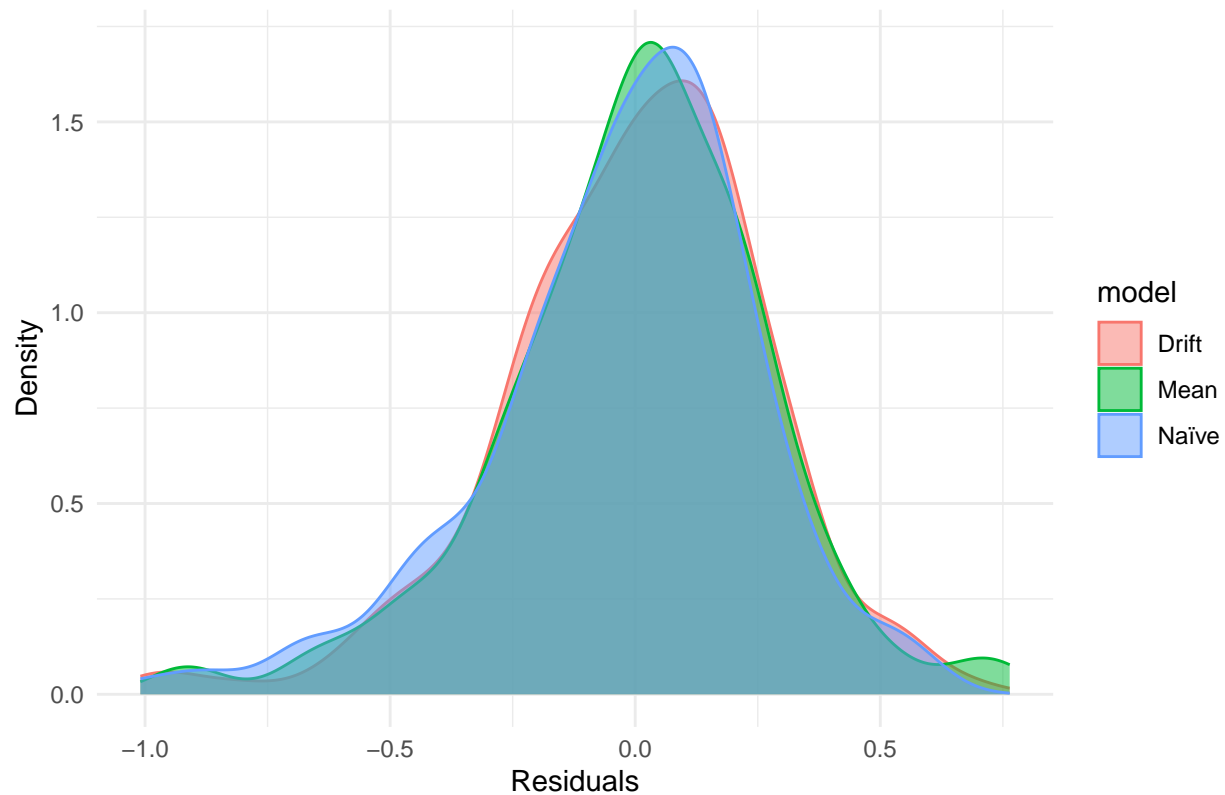
```



```
# Density plot of residuals
residuals |>
  ggplot(aes(x = .resid, fill = model, color = model)) +
  geom_density(alpha = 0.5) +
  theme_minimal() +
  labs(title = "Density Plot of Residuals for Forecasting Models", x = "Residuals", y = "Density")
```

```
## Warning: Removed 54 rows containing non-finite outside the scale range
## ('stat_density()').
```

Density Plot of Residuals for Forecasting Models



Especially the density plot shows that the residuals of the Naïve method are the closest to a normal distribution, which is a good sign for the accuracy of the model.

Actual Forecast with Naive Method

```
# Extend the forecast horizon to December 2024
last_year <- max(data$Year)
forecast_horizon <- 2024 - last_year

# Generate extended Naïve forecasts for each canton
extended_naive_fc <- data |>
  group_by(Kanton) |>
  model(naive_fc = NAIVE(Unemployment)) |>
  forecast(h = forecast_horizon)

# Extract the forecasted values
forecast_table <- extended_naive_fc |>
  as_tibble() |>
  select(Year, Kanton, .mean)

forecast_table
```

```
## # A tibble: 27 x 3
##   Year Kanton      .mean
```

```
##      <dbl> <chr>          <dbl>
## 1  2024 Aargau            2.23
## 2  2024 Appenzell A. Rh.  1.16
## 3  2024 Appenzell I. Rh.  0.544
## 4  2024 Basel-Landschaft  1.82
## 5  2024 Basel-Stadt       3.06
## 6  2024 Bern              1.42
## 7  2024 Freiburg         2.12
## 8  2024 Genf              3.78
## 9  2024 Glarus            1.30
## 10 2024 Graubünden        0.957
## # i 17 more rows
```

```
# TODO better display the values
```